



Slurm at UPPMAX

How to submit jobs with our
queueing system

Jessica Nettelblad
sysadmin at UPPMAX

Slurm at UPPMAX

- Intro
- **Queueing with Slurm**
 - **How to submit jobs**
- Testing
 - How to test your scripts before submission
- Monitoring
 - How to keep track on how the job is doing
- Examples

Free!

Watch!

Futurama S2 Ep.4
Fry and the Slurm factory



Simple Linux Utility
for Resource
Management

Open source!

<https://github.com/SchedMD/slurm>

Popular!

Used at many universities all
over the world



Submit jobs

Tell Slurm what the
job needs

Which script?

- Syntax:
 - `sbatch filename`

```
sbatch startjob.sh
```

Input after file name

- Syntax:
 - sbatch filename input arguments

```
sbatch startjob.sh input1 input2
```

Flags before filename

- Syntax:
 - sbatch flags filename input arguments

```
sbatch -A g2017016 -p core -n 1 -t 10:00 startjob.sh input1 input2
```



Flags in bash script

- Syntax:
 - sbatch filename input arguments

```
sbatch startjob.sh input1 input2
```

```
cat startjob.sh
```

```
#!/bin/bash
```

```
#SBATCH -A g2017016
```

```
#SBATCH -p core
```

```
#SBATCH -n 1
```

```
#SBATCH -t 10:00
```

```
start
```

```
Do this
```

```
Do that
```

```
end
```


Important flags


- **-A**
 - **Project name**
 - projinfo
 - **Example: -A g2017016**
- **-t**
 - **Run time**
 - **days-hours:minutes:seconds**
 - **Max: 10 days, 10-00:00:00**
 - **Example: -t 3-10:00:00**
- **-p**
 - **Partition**
 - One or more cores: core, devcore
 - One or more nodes: node, devel
 - **Example: -p core**
- **-n**
 - **Number of tasks (cores)**
 - Enough cores for threads:
Defined by the job script or program used
 - Enough cores for memory:
Test! Start low.
 - **Example: 1**

```
sbatch -A g2017016 -p core -n 1 -t 20:00 startjob.sh
```

Submit `startjob.sh` to Slurm, account it to `project g2017016`, let it run on `1 core` for `20 minutes`.

More flags

- **Job name**
 - -J testjob
- **--res=g2017016_WED**
 - Use reserved nodes
- Email notifications
 - --mail-type=FAIL
 - --mail-type=TIME_LIMIT_80
 - [--mail-user=jessica.nettelblad@it.uu.se](#)
- Output redirections
 - Default: work directory
 - --output=/proj/g2017016/nobackup/private/jessine/testjob/output/
- <https://slurm.schedmd.com/sbatch.html>
 - Most, but not all options are available at every center



Monitoring

Keep track of the job

- In queue
 - While running
 - When finished
-

Check queue

- ◉ **jobinfo**
 - ◉ Shows all jobs in queue. Modified squeue.
 - ◉ <https://slurm.schedmd.com/squeue.html>
- ◉ How many jobs are running?
 - ◉ `jobinfo | less`
 - ◉ Type q to exit
- ◉ When are my jobs estimated to start?
 - ◉ `jobinfo -u jessine`

Understanding the queue

- Priority
 - FAQ: <http://www.uppmax.uu.se/support/faq/running-jobs-faq/your-priority-in-the-waiting-job-queue/>
- Bonus jobs
 - Most projects has a quota of 2000 hours per 30 day
 - More than quota? Bonus job!
 - Can still run, no limits – but lower priority
 - How much have your project used? `projinfo`
- Dependency
 - Jobs won't start until another has finished successfully
 - Set with `--dependency`

Check progress

- ◉ **uppmix_jobstats**

- ◉ `less /sw/share/slurm/*/uppmix_jobstats/*/<job id>`
- ◉ Shows memory and core usage

- ◉ `scontrol show job <job id>`

- ◉ **ssh to the compute node**

- ◉ `top`
- ◉ `tail -f` (on result file)

Check finished job

- ◉ **slurm.out**
 - ◉ Check it for every job
 - ◉ Look for error messages
- ◉ **uppmix_jobstats**
- ◉ `finishedjobinfo -s today`

Cancel job

- ◉ Cancel all my jobs
 - ◉ `Scancel -u jessine`
- ◉ Cancel all my jobs named testjob
 - ◉ `scancel -j testjob`
- ◉ Cancel job with job id 4456
 - ◉ `scancel 4456`
 - ◉ Find job id? `jobinfo -u jessine`



Testing

Test using the

- interactive command
 - dev partition
 - fast lane
-

Testing in interactive mode

- ◉ `interactive` instead of `sbatch`
- ◉ No script needed
- ◉ `interactive -A g2017016 -t 15:00`
- ◉ Example:
 - ◉ A job script didn't work. I start an interactive job and submit line for line.

Testing in devel partition

- ◉ `-p devcore -n 4 -t 60`
 - ◉ Normal sbatch job: `sbatch flags jobscript input`
 - ◉ Run on four cores for 60 minutes
 - ◉ Note: Max one job submitted
- ◉ Job starts quickly!
- ◉ Example:
 - ◉ Before submitting the real job, I want to make sure the script and submit works as intended.

Testing in a fast lane

- `--qos=short`
 - Up to 15 minutes, up to four nodes
- `--qos=interact`
 - Up to 12 hours, up to one node
- Example:
 - My job is shorter than 15 minutes or 12 hours. I add a qos flag, and the job will get higher priority.
 - Note: Only one running job with qos – the rest will stay in queue.



Examples

Basic example

```
#!/bin/bash -l
#SBATCH -A g2016011
#SBATCH -p core
#SBATCH -n 1
#SBATCH -t 10:00:00
#SBATCH -J jour2
```

```
module load bioinfo-tools samtools/0.1.19 bwa
export SRCDIR=$HOME/baz/run3
```

```
cp $SRCDIR/foo.pl $SRCDIR/bar.txt $SNIC_TMP/
cd $SNIC_TMP
```

```
./foo.pl bar.txt
```

```
cp *.out $SRCDIR/out2
```

Basic example explained

- `#!/bin/bash`
 - starts the bash interpreter
- `#SBATCH -A g2016011`
 - `"#"` starts a comment that bash ignores
 - `"#SBATCH"` is a special signal to SLURM
 - `"-A"` specifies which account = project will be "charged".
- `#SBATCH -p core`
 - sets the partition to core, for jobs that uses less than one node.
- `#SBATCH -n 1`
 - requests one task = one core

Basic example explained

- `#SBATCH -t 10:00:00`
 - Time requested: 10 hours.
- `#SBATCH -J jour2`
 - `jour2` is the name for this job
 - mainly for your convenience
- `module load bioinfo-tools samtools/0.1.19 bwa`
 - `bioinfo-tools`, `samtools` version 0.1.19 and `bwa` is loaded.
 - can specify versions or use default (risky)
- `export SRCDIR=$HOME/baz/run3`
 - Environment variable `SRCDIR` is defined

Basic example explained

- `cp $SRCDIR/foo.pl $SRCDIR/bar.txt $SNIC_TMP/`
- `cd $SNIC_TMP`
 - Copy foo.pl and bar.txt to \$SNIC_TMP, then go there.
 - \$SNIC_TMP is a job specific directory on the compute nodes.
 - Recommended! Can be much faster than home.
- `./foo.pl bar.txt`
 - Actual script with code to do something.
 - Call one command, or a long list of actions with if-then, etc.
- `cp *.out $SRCDIR/out2`
 - \$SNIC_TMP is a temporary folder. It's deleted when job is finished.
 - Remember to copy back any results you need!

Group jobs

```
#!/bin/bash -l
#SBATCH -A g2016011
#SBATCH -p core
#SBATCH -n 4
#SBATCH -t 2-00:00:00
#SBATCH -J br_para_02
```

```
B=/sw/data/uppnex/igenomes/
S=Annotation/Genes/genes.gtf
```

```
grep AT2G48160 $B/Arabidopsis_thaliana/Ensembl/TAIR9/$S &
grep NM_030937 $B/Homo_sapiens/NCBI/GRCh38/$S &
grep NM_025773 $B/Mus_musculus/NCBI/GRCm38/$S &
grep XM_008758627.1 $B/Rattus_norvegicus/NCBI/Rnor_6.0/$S &
wait
```

Split jobs

```
#!/bin/bash

TOOL=z_tools/3
IGEN=/sw/data/uppnex/igenomes

cd $IGEN

for v in [[:upper:]]*/*/*
do
    echo $v
    cd $TOOL
    sbatch star_index.job $v
    cd $IGEN
    sleep 1
done
```

We're here to help!

- If you run into problems after this course? Just ask someone for help!
 - Check userguides and FAQ on uppmax.uu.se
 - Ask your colleagues
 - Ask UPPMAX support:
support@uppmax.uu.se



VOTED **#1** SOFT DRINK OF THE 31ST CENTURY!

